

Introduction to Statistical Learning – Final Project

Radha Gude

16337132

Part 1 :- Regression

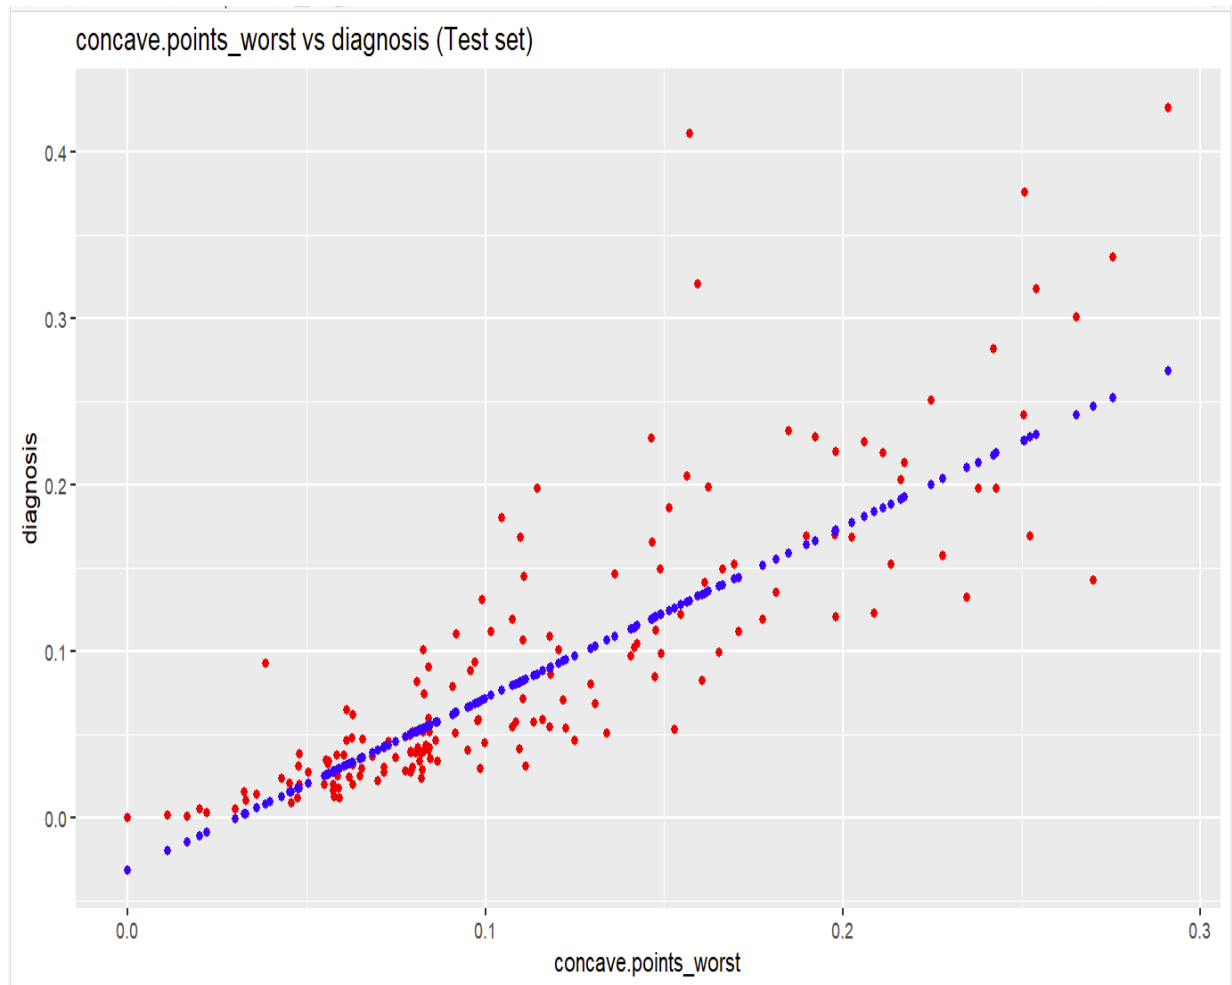
a) Linear Regression :

Code and Explanation :

```
#loading libraries
library('caTools')
library("ggplot2")
library(Metrics)
library(splines2)
#loading data set
df = read.csv('./Breastcancer.csv')
rsq <- function (x, y) cor(x, y) ^ 2
head(df)
cor(df)
# Splitting the dataset into the
# Training set and Test set
split = sample.split(df$id, SplitRatio = 0.7)
trainingset = subset(df, split == TRUE)
testset = subset(df, split == FALSE)
lm.fit <- lm( formula = concavity_mean ~ concave.points_worst, data = trainingset)
coef(lm.fit)
#testing testset
ypred = predict(lm.fit, newdata = testset)
summary(lm.fit)
modelPerformance = data.frame(
  RMSE = rmse(ypred, testset$concavity_mean),
  R2 = rsq(ypred, testset$concavity_mean)
)
#test results
print(modelPerformance)
#visualizing test set results

ggplot() +
  geom_point(aes(x= testset$concave.points_worst, y = testset$concavity_mean),
    colour = 'red') +
  geom_point(aes(x = testset$concave.points_worst,
    y = predict(lm.fit, newdata = testset)),
    colour = 'blue') +
  ggtitle('concave.points_worst vs diagnosis (Test set)') +
  xlab('concave.points_worst') + ylab('diagnosis')
```

Visualizing test results :



Summary :

```
Call:
lm(formula = concavity_mean ~ concave.points_worst, data = trainingset)

Residuals:
    Min       1Q   Median       3Q      Max
-0.094561 -0.024069 -0.003126  0.017883  0.226969

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)   -0.031666   0.003861  -8.201 3.35e-15 ***
concave.points_worst  1.029789   0.028836  35.711 < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.03753 on 396 degrees of freedom
Multiple R-squared:  0.7631,    Adjusted R-squared:  0.7625
F-statistic: 1275 on 1 and 396 DF,  p-value: < 2.2e-16

> modelPerformance = data.frame(
+   RMSE = rmse(ypred, testset$concavity_mean),
+   R2 = rsq(ypred, testset$concavity_mean)
+ )
> #test results
> print(modelPerformance)
      RMSE      R2
1 0.04701648 0.7100921
```

Observations :

Correlation between concavity_mean and concave.points_worst is nearly 0.67. So concavity_mean can be linearly dependent on concave.points_worst.

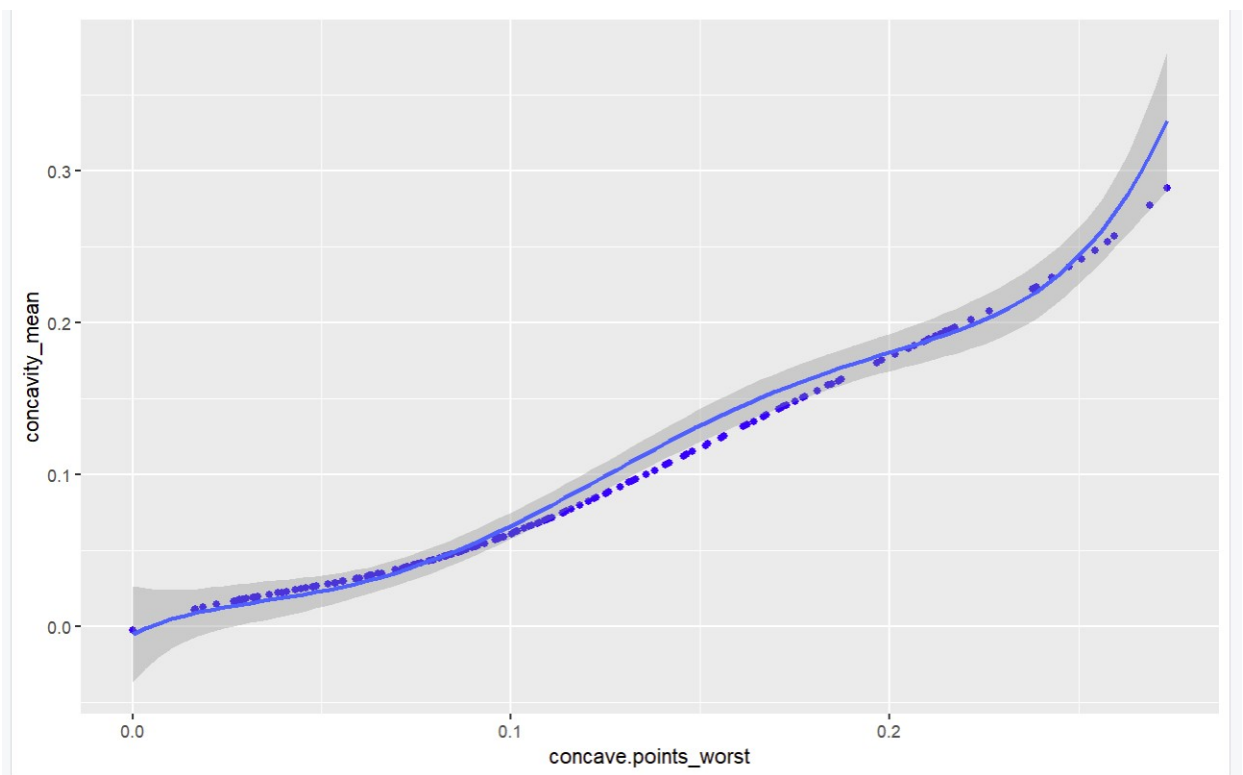
a) Polynomial Regression :

Code and Explanation :

```
#loading libraries
library('caTools')
library("ggplot2")
library(Metrics)
library(splines2)
#loading data set
df = read.csv('./Breastcancer.csv')
rsq <- function (x, y) cor(x, y) ^ 2
head(df)
cor(df)
# Splitting the dataset into the
# Training set and Test set
split = sample.split(df$id, SplitRatio = 0.7)
trainingset = subset(df, split == TRUE)
testset = subset(df, split == FALSE)
lm.fit <- lm( formula = concavity_mean ~ poly(concave.points_worst,5,raw = TRUE), data = trainingset)
coef(lm.fit)
#testing testset
ypred = predict(lm.fit, newdata = testset)
summary(lm.fit)
modelPerformance = data.frame(
  RMSE = rmse(ypred, testset$concavity_mean),
  R2 = rsq(ypred, testset$concavity_mean)
)
#test results
print(modelPerformance)
#visualizing test set results

ggplot(testset , aes(concave.points_worst,concavity_mean) , color = "red") +
geom_point(aes(x = testset$concave.points_worst,
               y = predict(lm.fit, newdata = testset)),
           colour = 'blue') +
  stat_smooth(method = lm, formula = y ~ poly(x, 5, raw = TRUE))|
```

Visualizing test results :



Summary :

```
Call:
lm(formula = concavity_mean ~ poly(concave.points_worst, 5, raw = TRUE),
    data = trainingset)

Residuals:
    Min       1Q   Median       3Q      Max
-0.137777 -0.018186 -0.005354  0.011599  0.284383

Coefficients:
                Estimate Std. Error t value Pr(>|t|)
(Intercept)      -2.301e-03  1.178e-02  -0.195   0.8453
poly(concave.points_worst, 5, raw = TRUE)1  1.081e+00  7.751e-01   1.394   0.1640
poly(concave.points_worst, 5, raw = TRUE)2 -1.915e+01  1.730e+01  -1.107   0.2689
poly(concave.points_worst, 5, raw = TRUE)3  2.330e+02  1.584e+02   1.471   0.1420
poly(concave.points_worst, 5, raw = TRUE)4 -1.015e+03  6.288e+02  -1.614   0.1073
poly(concave.points_worst, 5, raw = TRUE)5  1.529e+03  8.979e+02   1.702   0.0895 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.04066 on 392 degrees of freedom
Multiple R-squared:  0.7431,    Adjusted R-squared:  0.7398
F-statistic: 226.7 on 5 and 392 DF,  p-value: < 2.2e-16

> modelPerformance = data.frame(
+   RMSE = rmse(ypred, testset$concavity_mean),
+   R2 = rsq(ypred, testset$concavity_mean)
+ )
> #test results
> print(modelPerformance)
      RMSE      R2
1 0.0351015 0.8104029
```

Explanation :

Predicting concavity_mean with concave.points_worst using polynomial regression of order 5 yields results with R2 = 0.76959 which is really good.

C) Multi Linear Regression :

Code and Explanation:

```
#loading libraries
library('caTools')
library("ggplot2")
library(Metrics)
library(splines2)
#loading data set
df = read.csv('./Breastcancer.csv')
rsq <- function (x, y) cor(x, y) ^ 2
head(df)
cor(df)
# Splitting the dataset into the
# Training set and Test set
split = sample.split(df$id, splitRatio = 0.7)
trainingset = subset(df, split == TRUE)
testset = subset(df, split == FALSE)
lm.fit <- lm( formula = concavity_mean ~ .|, data = trainingset)
coef(lm.fit)
#testing testset
ypred = predict(lm.fit, newdata = testset)
summary(lm.fit)
modelPerformance = data.frame(
  RMSE = rmse(ypred, testset$concavity_mean),
  R2 = rsq(ypred, testset$concavity_mean)
)
#test results
print(modelPerformance)
```

Test Results :

```
> modelPerformance = data.frame(
+   RMSE = rmse(ypred, testset$concavity_mean),
+   R2 = rsq(ypred, testset$concavity_mean)
+ )
> #test results
> print(modelPerformance)
      RMSE      R2
1 0.01096864 0.9809185
> |
```

Summary :

Residuals:

Min	1Q	Median	3Q	Max
-0.032007	-0.004139	0.000380	0.004662	0.042229

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	5.368e-02	2.125e-02	2.526	0.01197 *
id	1.702e-12	4.389e-12	0.388	0.69843
diagnosis	3.135e-03	2.044e-03	1.534	0.12588
radius_mean	-4.260e-02	8.160e-03	-5.221	3.00e-07 ***
texture_mean	-5.579e-04	3.930e-04	-1.420	0.15660
perimeter_mean	6.153e-03	1.148e-03	5.357	1.50e-07 ***
area_mean	2.319e-05	2.551e-05	0.909	0.36397
smoothness_mean	-3.131e-01	9.633e-02	-3.250	0.00126 **
compactness_mean	1.342e-01	6.151e-02	2.181	0.02979 *
concave.points_mean	1.325e+00	6.690e-02	19.801	< 2e-16 ***
symmetry_mean	5.086e-02	3.535e-02	1.439	0.15109
fractal_dimension_mean	1.862e-01	2.718e-01	0.685	0.49373
radius_se	2.828e-02	1.420e-02	1.991	0.04720 *
texture_se	8.579e-04	1.669e-03	0.514	0.60745
perimeter_se	-5.407e-03	1.928e-03	-2.804	0.00532 **
area_se	1.418e-04	6.440e-05	2.202	0.02829 *
smoothness_se	3.465e-01	2.911e-01	1.190	0.23465
compactness_se	-2.971e-01	1.053e-01	-2.821	0.00505 **
concavity_se	5.683e-01	5.225e-02	10.875	< 2e-16 ***
concave.points_se	-5.548e-01	2.535e-01	-2.189	0.02924 *
symmetry_se	-6.459e-02	1.254e-01	-0.515	0.60682
fractal_dimension_se	6.847e-01	5.328e-01	1.285	0.19955
radius_worst	-3.565e-03	2.831e-03	-1.259	0.20872
texture_worst	4.433e-04	3.370e-04	1.316	0.18915
perimeter_worst	5.451e-04	2.874e-04	1.897	0.05863 .
area_worst	-1.152e-05	1.600e-05	-0.720	0.47229
smoothness_worst	1.614e-02	6.858e-02	0.235	0.81412
compactness_worst	-2.300e-02	1.807e-02	-1.273	0.20381
concavity_worst	1.539e-01	1.141e-02	13.497	< 2e-16 ***
concave.points_worst	-2.542e-01	4.149e-02	-6.126	2.33e-09 ***
concave.points_worst	-2.542e-01	4.149e-02	-6.126	2.33e-09 ***
symmetry_worst	-9.919e-03	2.347e-02	-0.423	0.67279
fractal_dimension_worst	-2.697e-01	1.170e-01	-2.305	0.02170 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.00946 on 366 degrees of freedom
Multiple R-squared: 0.9872, Adjusted R-squared: 0.9861
F-statistic: 907.3 on 31 and 366 DF, p-value: < 2.2e-16

D) Natural Cubic Spline:

Code and Explanation :

```
#loading libraries
library('caTools')
library("ggplot2")
library(Metrics)
library(splines2)
#loading data set
df = read.csv('./Breastcancer.csv')
rsq <- function (x, y) cor(x, y) ^ 2
head(df)
cor(df)
names(df)
# Splitting the dataset into the
# Training set and Test set
split = sample.split(df$id, SplitRatio = 0.7)
trainingset = subset(df, split == TRUE)
testset = subset(df, split == FALSE)
lm.fit <- lm( formula = concavity_mean ~ naturalSpline(concave.points_worst,df = 3), data = trainingset)
coef(lm.fit)
summary(lm.fit)
print(modelPerformance)
```

Test Results :

```
> print(modelPerformance)
              RMSE              R2
1 0.01096864 0.9809185
> |
```

Summary :

```
Call:
lm(formula = concavity_mean ~ naturalSpline(concave.points_worst,
  df = 3), data = trainingset)

Residuals:
    Min       1Q   Median       3Q      Max
-0.130495 -0.018504 -0.004347  0.012545  0.193191

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)    0.001446   0.007860   0.184   0.8541
naturalSpline(concave.points_worst, df = 3)1  0.042845   0.015660   2.736   0.0065 **
naturalSpline(concave.points_worst, df = 3)2 -0.236889   0.022490 -10.533 <2e-16 ***
naturalSpline(concave.points_worst, df = 3)3  0.899235   0.042006  21.407 <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.03577 on 394 degrees of freedom
Multiple R-squared:  0.7822,    Adjusted R-squared:  0.7805
F-statistic: 471.7 on 3 and 394 DF,  p-value: < 2.2e-16
```

Part 2 : Feature selection / Model Optimization Methods

Code :

```
library('caTools')
library("ggplot2")
library(Metrics)
library(leaps)
#loading data set
library(ISLR)
|
winedata = read.csv('./winequal.csv')

regfit.fwd <- regsubsets(quality ~ ., data = winequality,
                        nvmax = 11, method = "forward")
summary(regfit.fwd)
regfit.bwd <- regsubsets(quality ~ ., data = winequality,
                        nvmax = 11, method = "backward")
summary(regfit.bwd)

train <- sample(c(TRUE, FALSE), nrow(winedata),
               replace = TRUE)
test <- (!train)

regfit.best <- regsubsets(quality ~ .,
                        data = winequality[train, ], nvmax = 11)
reg.summaryfwd <- summary(regfit.fwd)
par(mfrow = c(1, 2))
plot(reg.summaryfwd $rss, xlab = "Number of Variables", ylab = "RSS", type = "l")
plot(reg.summaryfwd $adjr2, xlab = "Number of Variables", ylab = "Adjusted RSq", type = "l")
reg.summarybwd <- summary(regfit.bwd)
par(mfrow = c(1, 2))
plot(reg.summarybwd $rss, xlab = "Number of Variables", ylab = "RSS", type = "l")
plot(reg.summarybwd $adjr2, xlab = "Number of Variables", ylab = "Adjusted RSq", type = "l")
```

Forward:

Code:

```
library(ISLR2)
```

```
library(leaps)
```

```
Auto <- na.omit(Auto)
```

```
regfit.fwd <- regsubsets(alcohol ~ ., data = Auto,
```

```
                        nvmax = 10, method = "forward")
```

```
summary(regfit.fwd)
```

```
Call: regsubsets.formula(alcohol ~ ., data = Auto, nvmax = 10, method = "forward")
11 Variables (and intercept)
```

	Forced in	Forced out
fixed.acidity	FALSE	FALSE
volatile.acidity	FALSE	FALSE
citric.acid	FALSE	FALSE
residual.sugar	FALSE	FALSE
chlorides	FALSE	FALSE
free.sulfur.dioxide	FALSE	FALSE
total.sulfur.dioxide	FALSE	FALSE
density	FALSE	FALSE
pH	FALSE	FALSE
sulphates	FALSE	FALSE
quality	FALSE	FALSE

1 subsets of each size up to 10

Selection Algorithm: forward

	fixed.acidity	volatile.acidity	citric.acid	residual.sugar	chlorides
1 (1)	" "	" "	" "	" "	" "
2 (1)	" "	" "	" "	" "	" "
3 (1)	"*"	" "	" "	" "	" "
4 (1)	"*"	" "	" "	" "	" "
5 (1)	"*"	" "	" "	"*"	" "
6 (1)	"*"	" "	" "	"*"	" "
7 (1)	"*"	" "	" "	"*"	" "
8 (1)	"*"	" "	"*"	"*"	" "
9 (1)	"*"	"*"	"*"	"*"	" "
10 (1)	"*"	"*"	"*"	"*"	"*"

	free.sulfur.dioxide	total.sulfur.dioxide	density	pH	sulphates	quality
1 (1)	" "	" "	"*"	" "	" "	" "
2 (1)	" "	" "	"*"	" "	" "	"*"
3 (1)	" "	" "	"*"	" "	" "	"*"
4 (1)	" "	" "	"*"	"*"	" "	"*"
5 (1)	" "	" "	"*"	"*"	" "	"*"
6 (1)	" "	" "	"*"	"*"	"*"	"*"
7 (1)	"*"	" "	"*"	"*"	"*"	"*"
8 (1)	"*"	" "	"*"	"*"	"*"	"*"
9 (1)	"*"	" "	"*"	"*"	"*"	"*"
10 (1)	"*"	" "	"*"	"*"	"*"	"*"

> |

```

7:20 (Top Level) R
Console Terminal Background Jobs
R 4.3.0 ~/
subset selection object
Call: regsubsets.formula(alcohol ~ ., data = Auto, nvmax = 10, method = "backward")
11 Variables (and intercept)
              Forced in Forced out
fixed.acidity      FALSE      FALSE
volatile.acidity   FALSE      FALSE
citric.acid        FALSE      FALSE
residual.sugar     FALSE      FALSE
chlorides          FALSE      FALSE
free.sulfur.dioxide FALSE      FALSE
total.sulfur.dioxide FALSE      FALSE
density           FALSE      FALSE
pH               FALSE      FALSE
sulphates         FALSE      FALSE
quality          FALSE      FALSE
1 subsets of each size up to 10
Selection Algorithm: backward
fixed.acidity volatile.acidity citric.acid residual.sugar chlorides
1 ( 1 ) " " " " " " " "
2 ( 1 ) "*" " " " " " "
3 ( 1 ) "*" " " " " " "
4 ( 1 ) "*" " " " " "*" "
5 ( 1 ) "*" " " " " "*" "
6 ( 1 ) "*" " " " " "*" "
7 ( 1 ) "*" " " "*" "*" "
8 ( 1 ) "*" "*" "*" "*" "
9 ( 1 ) "*" "*" "*" "*" "
10 ( 1 ) "*" "*" "*" "*" "*"
free.sulfur.dioxide total.sulfur.dioxide density pH sulphates quality
1 ( 1 ) " " " " "*" " " " " "
2 ( 1 ) " " " " "*" " " " " "
3 ( 1 ) " " " " "*" "*" " " " "
4 ( 1 ) " " " " "*" "*" " " " "
5 ( 1 ) " " " " "*" "*" " " "*" "
6 ( 1 ) " " " " "*" "*" "*" "*" "
7 ( 1 ) " " " " "*" "*" "*" "*" "
8 ( 1 ) " " " " "*" "*" "*" "*" "
9 ( 1 ) " " "*" "*" "*" "*" "*" "
10 ( 1 ) " " "*" "*" "*" "*" "*" "

```

a) Forward Features

Code:

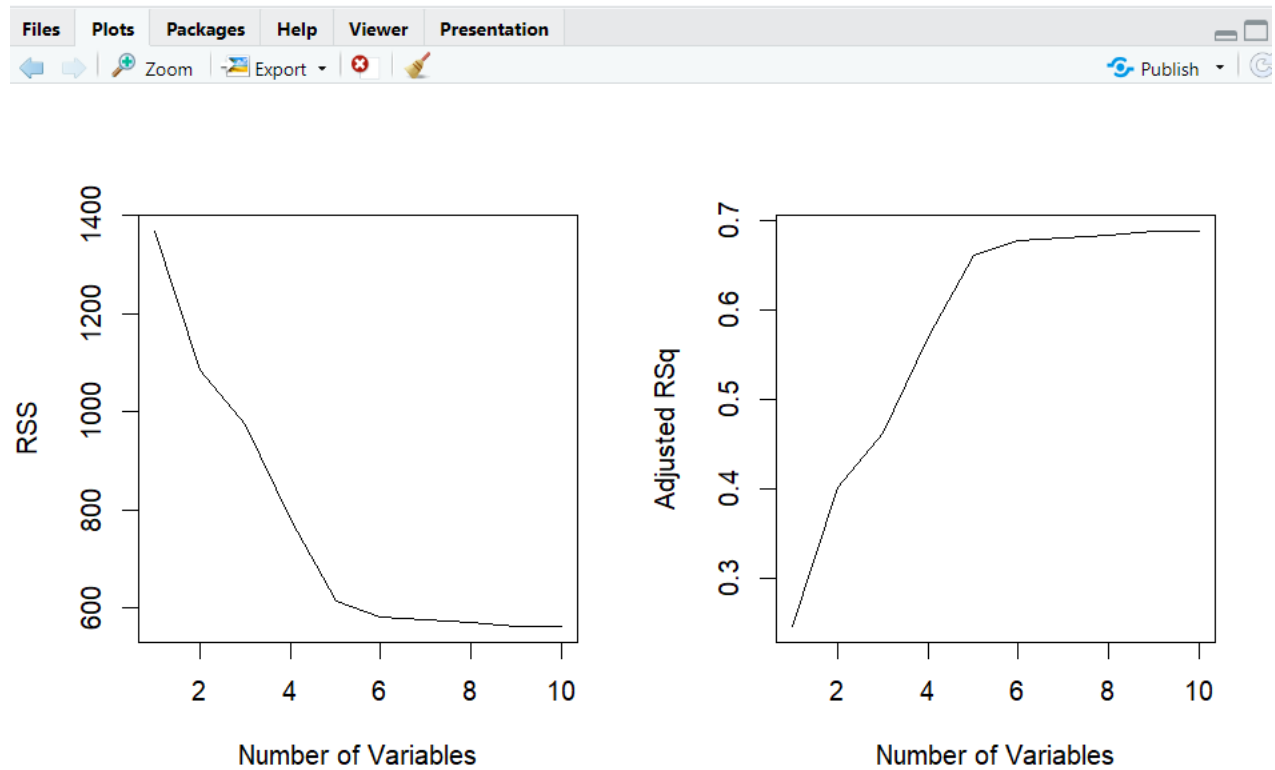
```
reg.summaryfwd <- summary(regfit.fwd)
```

```
par(mfrow = c(1, 2))
```

```
plot(reg.summaryfwd $rss, xlab = "Number of Variables", ylab = "RSS", type = "l")
```

```
plot(reg.summaryfwd $adjr2, xlab = "Number of Variables", ylab = "Adjusted RSq", type = "l")
```

Plots :



b) Backward Features

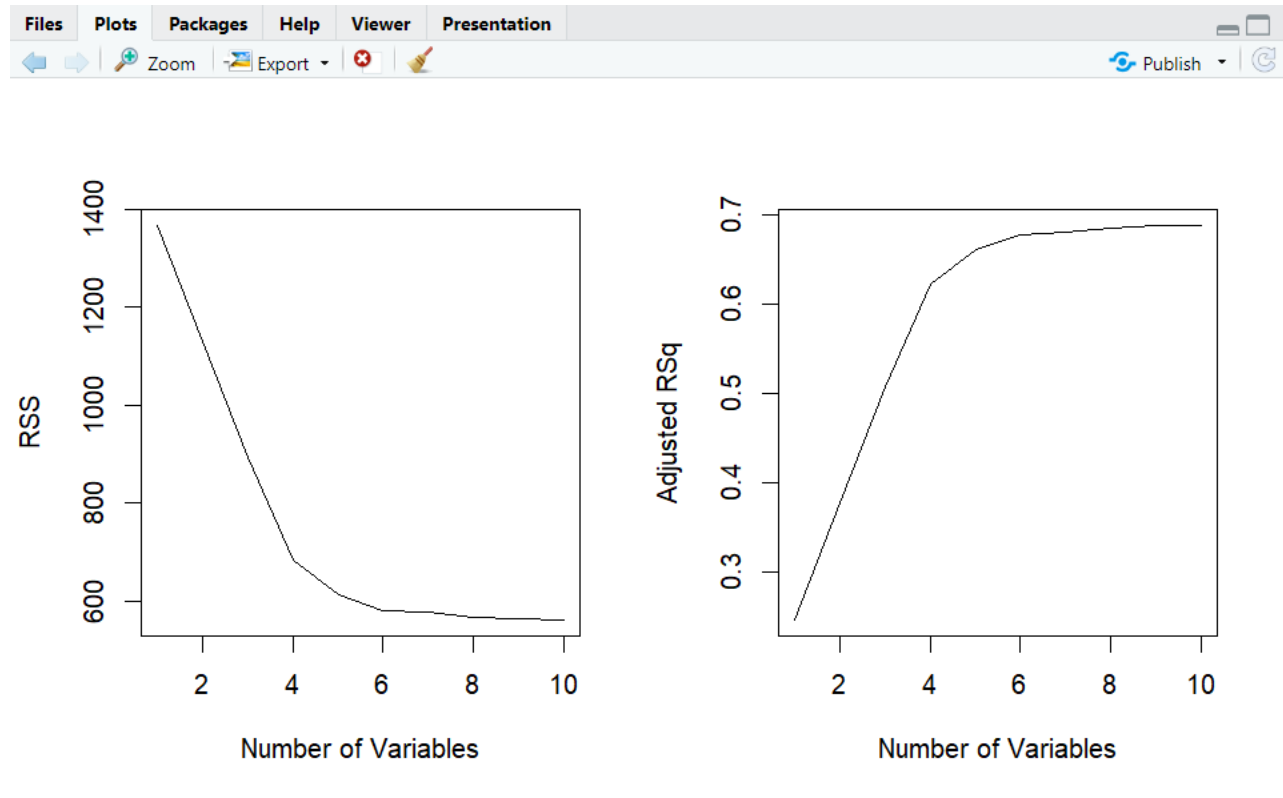
```
reg.summarybwd <- summary(regfit.bwd)
```

```
par(mfrow = c(1, 2))
```

```
plot(reg.summarybwd $rss, xlab = "Number of Variables", ylab = "RSS", type = "l")
```

```
plot(reg.summarybwd $adjr2, xlab = "Number of Variables", ylab = "Adjusted RSq", type = "l")
```

Plots:



Part 3 : Classification

A) Logistic Regression :

Code :

```
library(MASS)
df = read.csv('./heart_disease.csv')
head(df)
df<- df[-which(is.na(df$exang)), ]
split = sample.split(df$cp, SplitRatio = 0.7)
trainingset = subset(df, split == TRUE)
testset = subset(df, split == FALSE)
glm.fits <- glm(
  exang ~ sex+age+chol,
  data = trainingset, family = binomial
)
summary(glm.fits)
glm.probs <- predict(glm.fits, testset, type = "response")
glm.pred <- rep("FALSE", 260)
glm.pred[glm.probs > .5] = "TRUE"
table(glm.pred, testset$exang)
print(mean(glm.pred == testset$exang))
print(mean(glm.pred != testset$exang))
```

Result :

```
Call:
glm(formula = exang ~ sex + age + chol, family = binomial, data = trainingset)

Coefficients:
              Estimate Std. Error z value Pr(>|z|)
(Intercept) -3.4161459   0.6050171  -5.646 1.64e-08 ***
sexMale      0.7751653   0.2316384   3.346 0.000819 ***
age          0.0466468   0.0097158   4.801 1.58e-06 ***
chol        -0.0004241   0.0007932  -0.535 0.592925
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

    Null deviance: 789.02  on 584  degrees of freedom
Residual deviance: 748.65  on 581  degrees of freedom
(20 observations deleted due to missingness)
AIC: 756.65

Number of Fisher Scoring iterations: 4

> glm.probs <- predict(glm.fits, testset, type = "response")
> glm.pred <- rep("FALSE", 260)
> glm.pred[glm.probs > .5] = "TRUE"
> table(glm.pred, testset$exang)

glm.pred FALSE TRUE
FALSE    136    60
TRUE      28    36
> print(mean(glm.pred == testset$exang))
[1] 0.6615385
> print(mean(glm.pred != testset$exang))
[1] 0.3384615
> |
```

Explanation :

Based on above model we are able to predict exang with an 0.642 probability.

B) Linear Discriminant Analysis :

Code :

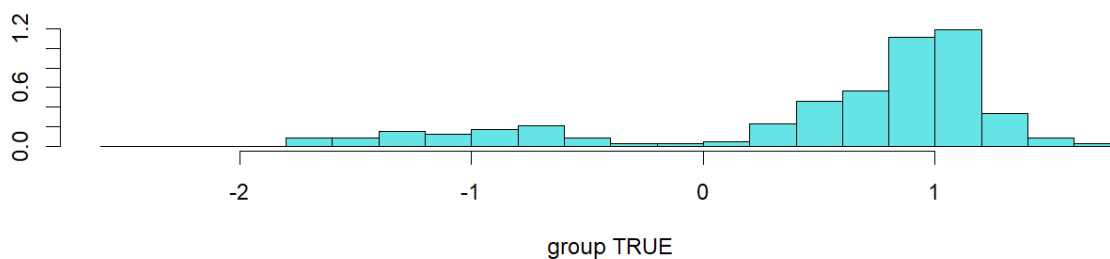
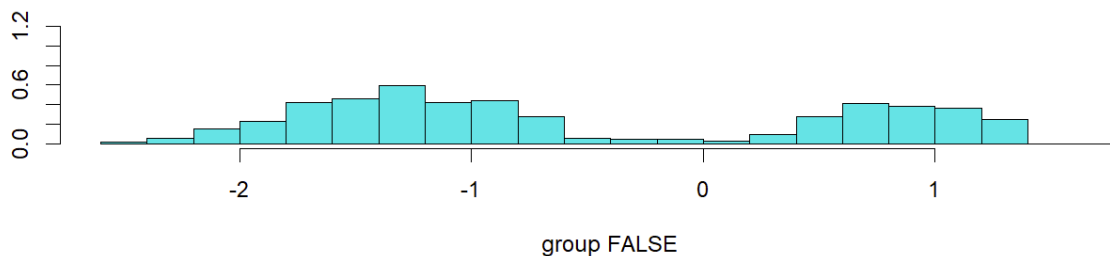
```
library(MASS)
df = read.csv('./heart_disease.csv')
head(df)
df<- df[-which(is.na(df$exang)), ]
split = sample.split(df$cp, SplitRatio = 0.7)
trainingset = subset(df, split == TRUE)
testset = subset(df, split == FALSE)
lda.fit <- lda(exang ~ cp + age + sex, data = trainingset)
plot(lda.fit)
lda.pred <- predict(lda.fit, testset)
lda.class <- lda.pred$class
table(lda.class, testset$exang)
mean(lda.class == testset$exang)
```

Results :

Confusion matrix is below:

```
lda.class FALSE TRUE
  FALSE   110   20
  TRUE    51   79
> mean(lda.class == testset$exang)
[1] 0.7269231
> |
```

Plots :



C) Tree classifier :

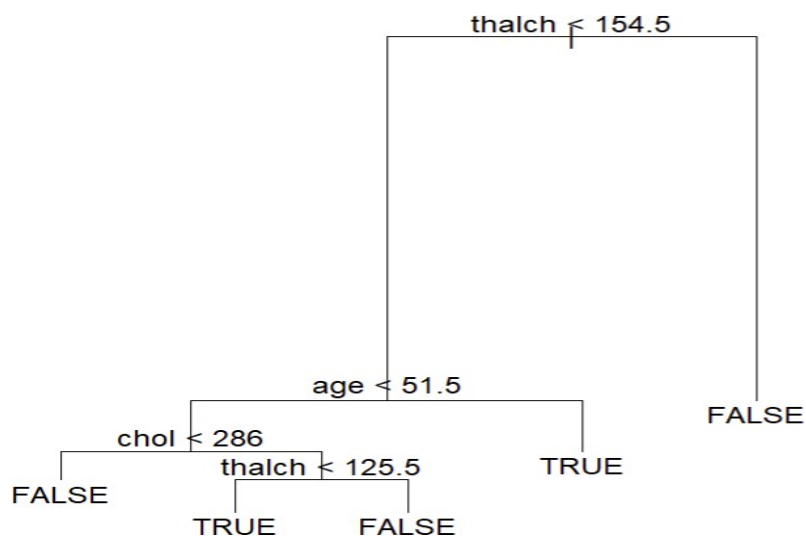
Code :

```
library(tree)
df = read.csv('./heart_disease.csv')
head(df)
df<- df[-which(is.na(df$exang)), ]
split = sample.split(df$cp, SplitRatio = 0.7)
trainingset = subset(df, split == TRUE)
testset = subset(df, split == FALSE)

tree.exang <- tree(as.factor(exang) ~ cp+age+sex+chol+fbs+thalch, trainingset)
summary(tree.exang)
plot(tree.exang)
text(tree.exang, pretty = 0)

tree.pred <- predict(tree.exang, testset,
                     type = "class" )
table(tree.pred, testset$exang)
```

Plots :



Confusion Matrix :

```
> table(tree.pred, testset$exang)

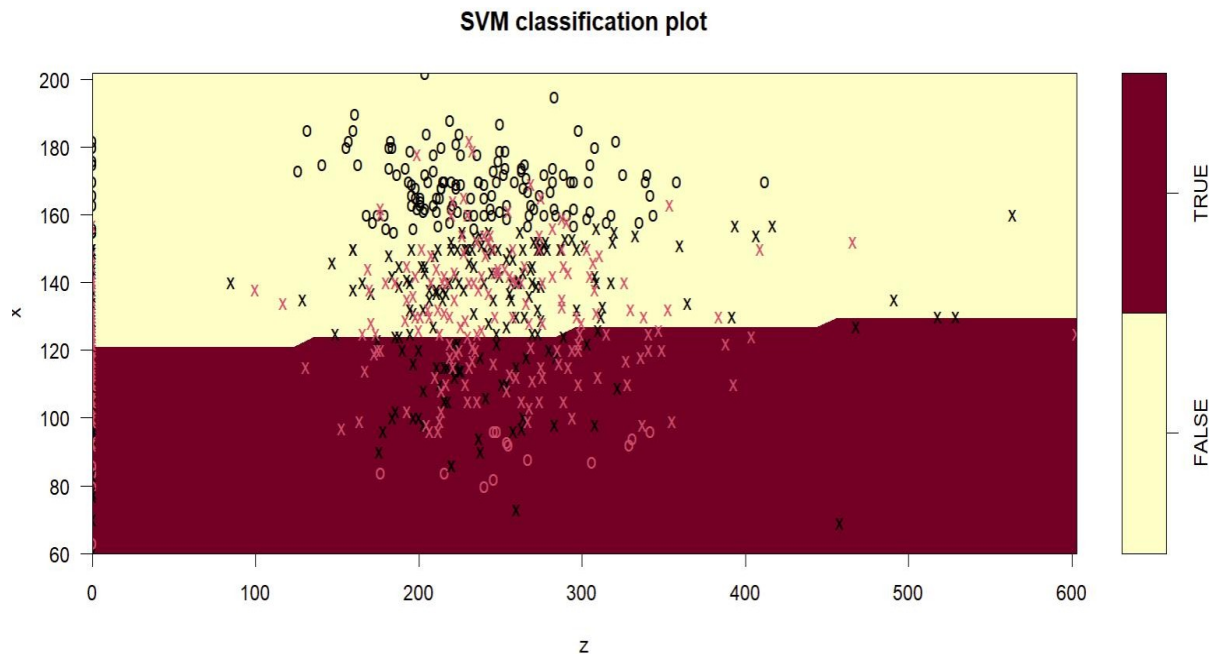
tree.pred FALSE TRUE
  FALSE    84    29
  TRUE     74    73
> |
```

3. SVM Linear classifier :

Code:

```
library('caTools')
library("ggplot2")
library(Metrics)
library(e1071)
df = read.csv('./heart_disease.csv')
df<- df[!which(is.na(df$exang)), ]
data = data.frame(x = df$thalch, y = as.factor(df$exang), z = df$chol);
split = sample.split(data$z, SplitRatio = 0.7)
trainingset = subset(data, split == TRUE)
testset = subset(data, split == FALSE)
svmfit <- svm(y ~ x+z, data = trainingset, kernel = "linear", cost = 1, scale = TRUE)
#generate a plot for submission in the next step.
summary(svmfit)
plot(svmfit, trainingset)
#using the built in tune function, we adjust the value of "cost" a few times (7 values) and see the results for a 10-Fold Cross Validation
tune.out <- tune(svm, y ~ x+z, data = data, kernel = "linear",
               ranges = list(cost = c(0.001, 0.01, 0.1, 1, 5, 10, 100)))
#we output the "best" performance model from the k-fold above
summary(tune.out)
#we copy that model and use summary again to show the parameters of that training
bestmod <- tune.out$best.model
summary(bestmod)
```

SVM Plot :



Best Model Summary :

```
- best performance: 0.3319198
- Detailed performance results:
  cost      error dispersion
1 1e-03 0.3948920 0.04154202
2 1e-02 0.3781816 0.03561521
3 1e-01 0.3319198 0.05050435
4 1e+00 0.3378585 0.04659565
5 5e+00 0.3390490 0.04762182
6 1e+01 0.3390490 0.04762182
7 1e+02 0.3390490 0.04762182

> #we copy that model and use summary again to show the parameters of that training
> bestmod <- tune.out$best.model
> summary(bestmod)

Call:
best.tune(METHOD = svm, train.x = y ~ x + z, data = data, ranges = list(cost = c(0.001, 0.01,
0.1, 1, 5, 10, 100)), kernel = "linear")

Parameters:
  SVM-Type:  C-classification
  SVM-Kernel: linear
    cost:  0.1

Number of Support Vectors:  620

( 310 310 )

Number of Classes:  2

Levels:
  FALSE TRUE
```

Explanation :

With this SVM classification I tried to predict exang based on chol and thalch. I was able to generate classifier with probability of 0.62.

Part 4 : Models

A) (5 points) A friend is starting a company and wants your help to see if they can figure out what factors most closely relate to the relative level of success for key competitors. They have gathered a few factors about each company such as total inventory, number of employees, annual operation budget and total profits. What method might you use to help your friend determine if their business model might be a success? Why did you choose this model?

Ans) Total inventory, headcount, annual operational budget, and total revenue are the four ways to identify anything. Every variable seems to have a linear relationship. To assess the likelihood that his business plan will be successful, I would suggest carrying out a multiple linear regression study.

B) (5 points) An advertisement firm has hired you to help them optimize their mailing list. They currently are looking to promote their client's store by sending packages of coupons to select areas. We want to know which postal codes the company should mail to for maximum impact (shoppers come to the store with coupons). They currently have some survey data randomly sampled from homes in the area indicating how likely they were to shop at the client's location. What method might you try first to generate the mailing map? Why?

Ans) We can use logistic regression to ascertain the likelihood that a customer will shop at the client's location. Based on overall results, we can pinpoint the particular regions where mailing will have the greatest impact.

C) (5 points) A large company has been collecting data about their customers preferences for many years. They've hired you to help them transform the millions of samples and thousands of search and behavior features into a set of simplified features they can use to build a model which provides suggestions to their customers for future services. What method might you suggest first? Why?

Ans) We must transform the enormous amount of data we have into traits that are simpler to comprehend. We can use PCA to get a linear combination of variables. The finest regression algorithms can then be chosen in order to acquire additional insights.

d)(5 points) A company that specializes in shipping fruit to grocery stores wants to save money by sorting out bad fruit from good fruit before it goes on the truck. They have presented you with a device that can measure features like weight, color, size, and look for possible bad spots. Each of these measurements is imprecise, and there is significant overlap between the classes for most of the features. What supervised learning methods might you try? Why?

Ans) It is difficult to categorise this since, as was already mentioned, there is a lot of overlap between classes. I would have thought about multidimensional SVM if the characteristics were easily separable. However, it seems that the random forest classifier can more accurately handle these overlaps.